

Part 1

A1.

Research question: What is the combined effect of a customer's geographic location, service outages, and monthly charges on their overall tenure with telecommunications companies?

A2.

The goal of my analysis is to determine how customer's factors influence their tenure. Utilizing multiple linear regression should permit this investigation, and the ability to provide recommendations to the telecommunication companies.

Part 2

B1.

Multiple linear regression can only be applied if all technical conditions are met. These assumptions are critical for the accuracy and reliability of the model's parameter estimates and predictions. The four assumptions of a multiple linear regression model include:

1. **Linearity:** The relationship between the independent and dependent variables follows a linear model. The changes in the dependent variables are proportional to changes in the independent variables.
2. **Independence:** The observations in the analysis are independent of each other, meaning the values of each dependent variable do not influence observations.
3. **Normality:** The differences between the observed values and the predicted values, residuals, are normally distributed. This condition ensures that the p-value and confidence intervals estimates are an accurate reflection of the population values.
4. **Equal variability:** The spread of the residuals should be approximately the same for all values of the independent variables.

B2.

Python is the selected programming language because the data set did not require detailed visualizations over periods of time, and it did not require intense statistical technique to clean data. Python is a much more general, robust program that does not require specific libraries and syntax, which makes it easier and more time effective to achieve our goal of conducting multiple linear regression. Furthermore, pandas, numpy, missingno, seaborn, matplotlib, scikitlearn, fancy impute, and statsmodels were all used to conduct various phases of the analysis. The various phases of the analysis for this report included data cleaning, data transformation, feature selection, model construction, and statistical calculations. In comparison to R, I would have had to use specific syntax for each library installed, and this could have unnecessarily increased the complexity of my investigation.

B3.

Multiple linear regression technique allows you to model the relationship between the continuous dependent variable and multiple independent variables or predictors simultaneously, which allows us to better understand real world problems that involve how multiple factors jointly affect a decision. This is the appropriate technique for analyzing the research question because the target variable, the customer's tenure, is a continuous variable. Also, it allows me to investigate how multiple independent variables such as customer's geographic location, service outages, and monthly charges influence the continuous dependent variable, the customer's tenure.

Multiple linear regression provides quantitative insights into the strength and direction of the relationship between the independent variables and the dependent variable. This helps to estimate the coefficients of the independent variables to indicate the magnitude of their impact on the dependent variable, customer tenure. Also, this technique assumes a linear relationship between the independent variables and the dependent variable. This aligns with the idea that changes in predictor or independent variables should have a linear effect on the dependent variable, customer tenure. Additionally, multiple linear regression offers the ability to assess the statistical significance of the independent variables to determine which factors are most influential in predicting customer tenure.

In conclusion, multiple linear regression is the appropriate technique for analyzing the research question because it accommodates the continuous dependent variable and allows for the investigation of multiple independent variables at once, provides insights into the relationships, and offers statistical tests to assess the significance of these relationships.

In text citations:

("Technical conditions for linear regression", n.d.)

Part 3

C1.

The data cleaning process entails investigating missing and duplicate values to decide how to impute missing values. The goal of this process is to gain insight into missing variables' values for effective data imputation. Clean data is critical for preparing the dataset to conduct multiple linear regression and provide meaningful recommendations to answer the research question.

Data Cleaning Plan:

Step 1: Exploratory Data Analysis

1. Loaded CSV file into Jupyter notebook using pandas library and converted the CSV file into a Data Frame.
2. Used the methods, 'describe()', 'info()', and 'dtypes()' to observe the summary statistics, data types, and nullity, the data types of the data frame.
3. Identified if there was duplicated information using the 'duplicated()' method.
4. Visualized missing data across the entire data frame using missingno and matplotlib libraries, using the 'bar()' and 'show()' methods.
5. Created a correlation matrix using the 'corr()' method, and utilized the correlation matrix to create a heatmap using the seaborn library to understand the correlation between variables missing values.
6. Determined missingness type(s) if applicable.

Step 2: Clean Data

1. Upon completion of investigation, the variables below had missing values:
 - a. Children- 7505 (numerical)
 - b. Income- 7510 (numerical)
 - c. Techie- 7523 (non-numerical)

- d. Age- 7525 (numerical)
 - e. Phone- 8974 (non-numerical)
 - f. Bandwidth GB Year- 8979 (numerical)
 - g. Tech Support- 9009 (non-numerical)
 - h. Tenure- 9069 (numerical)
2. By analyzing the summary statistics, I noticed that the population had a minimum value of zero. Since it is highly unlikely population is zero, I replaced all population records with values equal to zero with NaN values utilizing the numpy library and the 'nan' and 'isnan' methods. This allows Python to recognize missing values properly prior to imputation. The missing Population values were imputed using the K-Nearest Neighbor technique.
 3. The numerical values appear to be missing at random because the missing values could probably be predicted by other observed values in the data set. For example, missing number of children values could be relative to age and/or marital status (e.g., if a customer is an unmarried young adult, they are likely to be within the age range 18-30, and not list the number of children that they have since it might not be applicable). To treat numerical missing values, I decided to use the K-Nearest Neighbor imputation technique because missing values can likely be predicted based on similar observed variables and values.
 4. The non-numerical values also appear to be missing at random because the missing values could also probably be predicted by other observed values in the data set. For example, missing Techie values could also be relative to age (e.g., if a customer is between ages 50-90, it is unlikely that they would be considered a Techie). Likewise, missing values for Phone and Tech Support could likely be related to age due to the fast advancement of technology overtime. To impute non-numerical variables' values, I decided to use the Simple Imputation mode technique, and used frequency to fill in missing data. Furthermore, to validate this imputation technique, I analyzed the clean data to ensure that it fit the well-considered estimate, and the clean data portrayed that the group that are not techie's lies between ages 50-90.

Code for C1:

See code/script attached and below:

Please see Churn Data.ipynb attached for this code.

```
import pandas as pd
import numpy as np
import missingno as msno
churn_raw_data = pd.read_csv ('/Users/jasminemoniquecooper/Downloads/churn_raw_data.csv')
pd.set_option('display.max_columns', None)
churn_raw_data.head(10)

churn_raw_data.dtypes

churn_raw_data.describe()

churn_duplicates = churn_raw_data.duplicated()
churn_raw_data[churn_duplicates]

churn_raw_data.isna().sum()
```

```

#overall summary of missing data in the data frame

#visualize missingness
column_order = churn_raw_data.isnull().sum().sort_values().index
msno.bar(churn_raw_data[column_order])
plt.show()

import seaborn as sns
import matplotlib.pyplot as plt

# Generate a correlation matrix
correlation_matrix = churn_raw_data.corr()

# Set the figure size to make the heatmap larger
plt.figure(figsize=(12, 10))

# Create a heat map using seaborn with customizations
heatmap = sns.heatmap(correlation_matrix, annot=True, cmap='coolwarm', fmt=".2f", linewidths=0.5,
annot_kws={"size": 8})

# Increase the font size of the color bar (optional)
cbar = heatmap.collections[0].colorbar
cbar.ax.tick_params(labelsize=12)

# Increase the font size of the annotations
for text in heatmap.texts:
    text.set_size(8)

# Display the heat map
plt.show()

churn_raw_data.Population[churn_raw_data.Population == 0].count()

#determine if missing income values are related to employment

# Select the rows where income is missing
missing_income = churn_raw_data[churn_raw_data['Income'].isnull()]

# Group the missing income data by employment status
grouped = missing_income.groupby(['Employment'])

# Count the number of missing income values for each employment status
missing_count = grouped.size()

# Calculate the percentage of missing income values for each employment status
percentage_missing = missing_count / churn_raw_data.groupby(['Employment']).size() * 100

```

```

# Display the results
print(percentage_missing)
#determine if missing income values are related to employment and age

# Select the rows where income is missing
missing_income = churn_raw_data[churn_raw_data['Income'].isnull()]

# Group the missing income data by employment status
grouped = missing_income.groupby(['Employment', pd.cut(missing_income['Age'], bins=[0, 18, 30, 50,
np.inf])])

# Count the number of missing income values for each employment status
missing_count = grouped.size()

# Calculate the percentage of missing income values for each employment status
percentage_missing = missing_count / churn_raw_data.groupby(['Employment',
pd.cut(churn_raw_data['Age'], bins=[0, 18, 30, 50, np.inf])]).size() * 100

# Display the results
print(percentage_missing)

#part time employment between ages 0-18 is missing data the most

import pandas as pd

# Select the rows where the 'Techie' column is null
missing_techie = churn_raw_data[churn_raw_data['Techie'].isnull()]

# Define age group bins
age_bins = [18, 30, 50, float('inf')]

# Group the missing 'Techie' data by age
grouped = missing_techie.groupby(pd.cut(missing_techie['Age'], bins=age_bins))

# Count the number of missing 'Techie' values for each age group
missing_count = grouped.size()

# Calculate the percentage of missing 'Techie' values for each age group
total_count_by_age_group = churn_raw_data.groupby(pd.cut(churn_raw_data['Age'],
bins=age_bins)).size()
percentage_missing = (missing_count / total_count_by_age_group) * 100

# Display the results
print(percentage_missing)

import pandas as pd

```

```

# Select the rows where 'number of children' is missing
missing_children = churn_raw_data[churn_raw_data['Children'].isnull()]

# Define age group bins
age_bins = [18, 30, 50, float('inf')] # Adjust the bins as needed

# Group the missing 'number of children' data by marital status and age
grouped = missing_children.groupby(['Marital', pd.cut(missing_children['Age'], bins=age_bins)])

# Count the number of missing 'number of children' values for each marital status and age group
missing_count = grouped.size()

# Calculate the percentage of missing 'number of children' values for each marital status and age group
total_count_by_group = churn_raw_data.groupby(['Marital', pd.cut(churn_raw_data['Age'],
bins=age_bins)]).size()
percentage_missing = (missing_count / total_count_by_group) * 100

# Display the results
print(percentage_missing)

import numpy as np
churn_raw_data.loc[churn_raw_data.Population == 0, 'Population'] = np.nan
churn_raw_data.Population[np.isnan(churn_raw_data.Population)]

pip install fancyimpute
from fancyimpute import KNN

columns_to_impute = ['Children', 'Income', 'Age', 'Bandwidth_GB_Year', 'Population', 'Tenure']

# Create an instance of the KNN imputer
knn_imputer = KNN()

# Get the column indices of the columns to impute
columns_to_impute_indices = [churn_raw_data.columns.get_loc(col) for col in columns_to_impute]

# Perform imputation on the selected columns
imputed_values = knn_imputer.fit_transform(churn_raw_data.iloc[:, columns_to_impute_indices])

# Assign the imputed values back to the original dataset
churn_raw_data.iloc[:, columns_to_impute_indices] = imputed_values

# Verify if the imputations are in the original dataset
print(churn_raw_data.head())

from sklearn.impute import SimpleImputer

column_to_impute_two = ['Techie', 'Phone', 'TechSupport']

```

```

# Create the SimpleImputer object with strategy='most_frequent'
imputer = SimpleImputer(strategy='most_frequent')

# Fit and transform the categorical variables using the imputer
churn_raw_data[column_to_impute_two] =
imputer.fit_transform(churn_raw_data[column_to_impute_two])

import pandas as pd
# Define age group bins
age_bins = [18, 30, 50, float('inf')]

# Group the data by age
grouped = churn_raw_data.groupby(pd.cut(churn_raw_data['Age'], bins=age_bins))

# Count the number of 'Techie' values for each age group
tech_count = grouped['Techie'].value_counts().unstack(fill_value=0)

# Calculate the percentage of 'Techie' values for each age group
tech_percentage = (tech_count / tech_count.sum(axis=1).values[:, None]) * 100

# Display the results
print(tech_percentage)

#data clean
import matplotlib.pyplot as plt
column_order = churn_raw_data.isnull().sum().sort_values().index
msno.bar(churn_raw_data[column_order])
plt.show()

cleaned_data_two = churn_raw_data
new_name_two = 'churn_cleaned_data'
new_cleaned_data_two = cleaned_data_two.copy()
new_cleaned_data_two.name = new_name_two
new_cleaned_data_two.to_csv('new_cleaned_data_two.csv', index=False)
churn_raw_data.to_csv(r'/Users/jasminemoniquecooper/Downloads/new_cleaned_data_two.csv')

```

C2.

For the dependent variable, tenure, I have 10,000 observations to examine the summary statistics:

- The mean indicates that on average customers stay with telecommunications companies for approximately 34 days
- The standard deviation of 25 days shows that individual customer tenures can vary from the mean by plus or minus 25 days
- The minimum tenure observed is 1 day, while the maximum tenure is approximately 71 days which is indicative of the range of customer durations
- The 25th percentile is around 8 days, meaning that 25% of customers stay for 8 days or less
- The 75th percentile is around 60 days, meaning that 75% of customers stay for 60 days or less
- The median, which is around 36 days, meaning that 50% of customers stay for more than 36 days, and 50% of customers stay for less than 36 days

	Tenure
count	10000.000000
mean	34.624193
std	25.903916
min	1.000259
25%	8.197844
50%	36.852175
75%	60.457955
max	71.999280

The proximity of the mean and median suggests that the distribution may approximate a bell curve or normal distribution.

For the independent variables, Outage_sec_perweek, MonthlyCharge, and Area I have 10,000 observations to examine the summary statistics:

Independent variable 1: Outage_sec_perweek – average number of seconds per week of system outages in the customer's neighborhood

- The mean indicates that on average customers experience an 11 second average outage in their neighborhood per week
- The standard deviation of 7 seconds shows that average outages per week in the customer's neighborhood can vary from the mean by plus or minus 7 seconds
- The minimum seconds of outages per week is negative, which is not a valid measure of time, so I will say that the minimum amount of time a customer on average experiences an outage in their neighborhood is 0 seconds. In comparison, the maximum amount of time a customer may on average experience an outage in their neighborhood is 47 seconds.
- The 25th percentile is around 8 seconds, meaning that 25% of customers on average experience an outage of 8 seconds per week or less
- The 75th percentile is around 12 seconds, meaning that 75% of customers on average experience an outage of 12 seconds per week or less
- The median, which is around 10 seconds, meaning that 50% of customers on average experience an outage for more than 10 seconds per week, and 50% of customers experience an outage, on average, for less than 10 seconds per week

The proximity of the mean and median suggests that the distribution may approximate a bell curve or normal distribution.

Independent variable 2: MonthlyCharge – The amount charged to the customer monthly

- The mean indicates that on average customers are charged \$174 monthly
- The standard deviation of \$43 shows that customer's monthly charges can vary from the mean by plus or minus \$43

- The minimum amount a customer is charged monthly is \$77, while the maximum monthly charge to a customer is \$315, which is indicative of the range of customer monthly charges
- The 25th percentile is \$141, meaning that 25% of customers receive a monthly charge of \$141 or less
- The 75th percentile is \$203, meaning that 75% of customers receive a monthly charge of \$203 or less
- The median, which is \$169, meaning that 50% of customers receive a monthly charge more than \$169, and 50% of customers receive a monthly charge less than \$169

The proximity of the mean and median suggests that the distribution may approximate a bell curve or normal distribution.

Independent variable 3: Area – area type (rural, urban, or suburban)

Since the area is a non-numerical, categorical independent variable then I can evaluate summary statistics by examining the frequency of occurrence for each area type.

Frequency of Occurrence for Each Area Type:

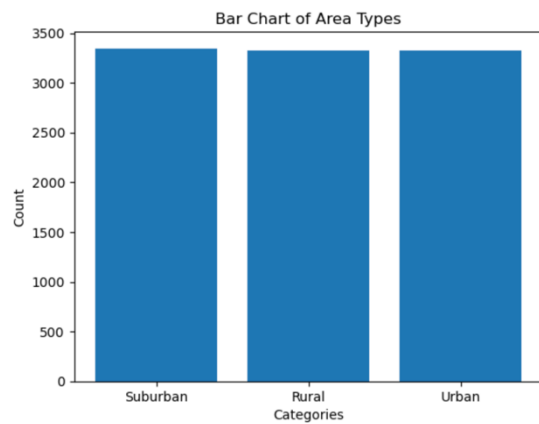
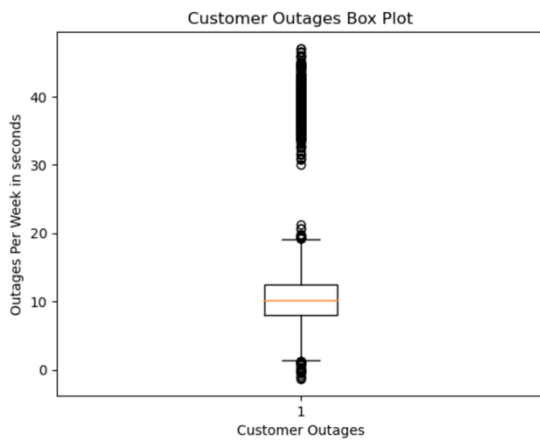
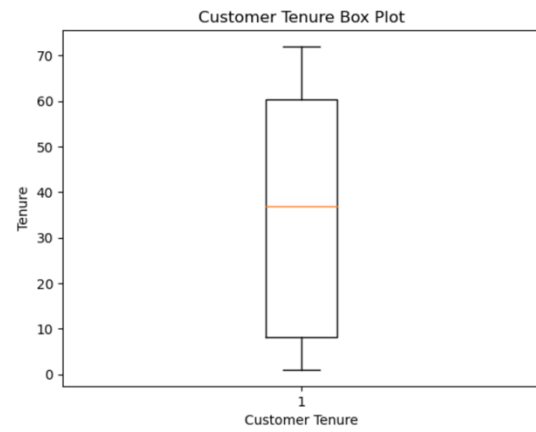
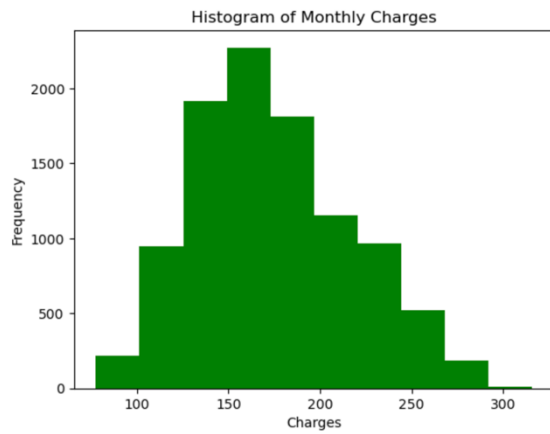
- Rural – count of observations in the rural area
- Urban – count of observations in the urban area
- Suburban – count of observations in the suburban area

Based on the output below, it appears that the 10,000 observations are evenly distributed between all 3 area types: rural, urban, and suburban. The balance in the distribution of area types indicates that the dataset provides a representative sample, which is critical for drawing meaningful conclusions in our research.

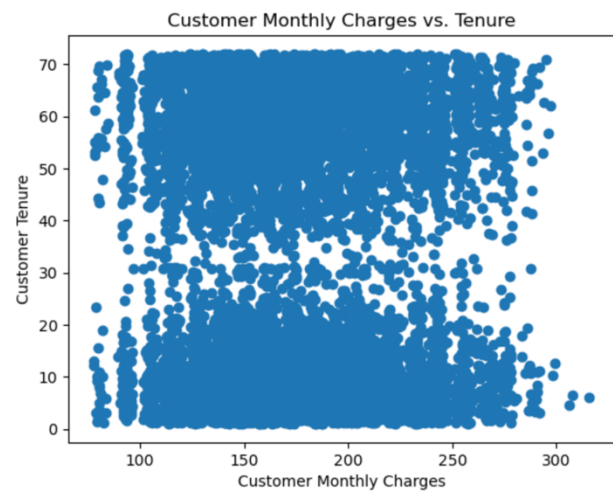
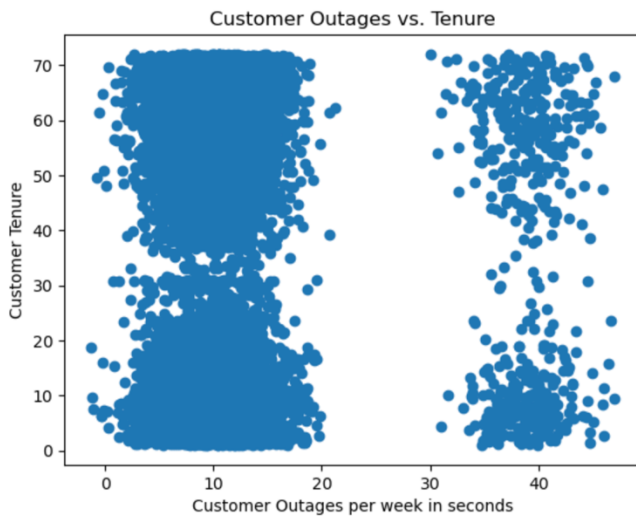
	MonthlyCharge	Outage_sec_perweek	
count	10000.000000	10000.000000	Mode (most frequent 'Area'): Suburban
mean	174.076305	11.452955	Frequency counts of 'Area':
std	43.335473	7.025921	Suburban 3346
min	77.505230	-1.348571	Rural 3327
25%	141.071078	8.054362	Urban 3327
50%	169.915400	10.202896	Name: Area, dtype: int64
75%	203.777441	12.487644	Percentage distribution of 'Area':
max	315.878600	47.049280	Suburban 33.46
			Rural 33.27
			Urban 33.27
			Name: Area, dtype: float64

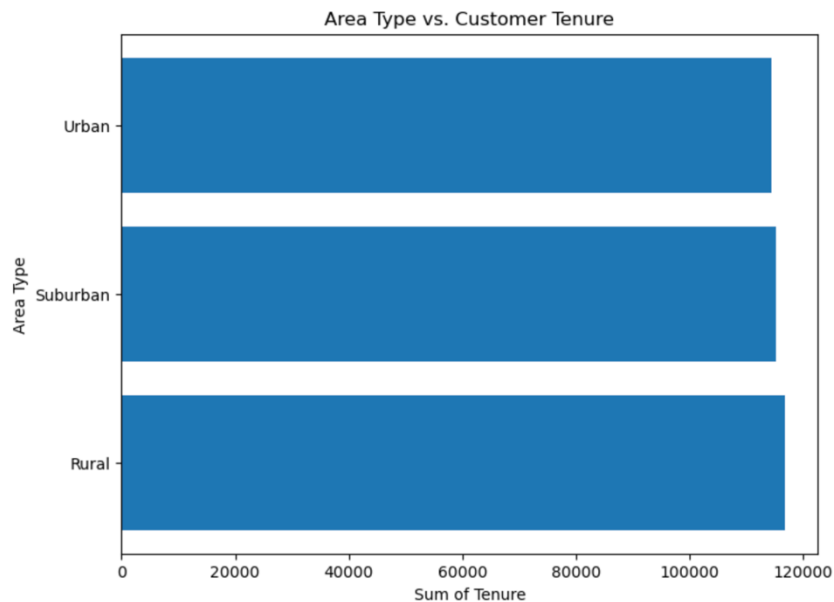
C3.

Univariate visualizations:



Bivariate visualizations:





C4.

Converting columns from data type object to data type category was necessary because there are only a few different values to be selected from each variable and it helps to save memory and ensures that Python recognizes these variables as categorical variables in other libraries and visualizations. Specifically, converting the Area variable allows me to properly calculate the frequency of each area type and portray accurate univariate and bivariate visualizations. Furthermore, to effectively conduct feature selection for the reduced model, I employed one hot encoding, which created three separate binary variables for the Area variable.

List of columns converted from data type object to data type category:

- Employment, Marital, Gender, Churn, Techie, Contract, Post_modem, Tablet, InternetService, Phone, Multiple, OnlineSecurity, OnlineBackup, DeviceProtection, TechSupport, StreamingTV, StreamingMovies, PaperlessBilling, PaymentMethod, and Area

Code for C4 (changed variable names as needed):

```
churn_clean_data["Area"] = churn_clean_data["Area"].astype('category')
churn_clean_data["Area"].dtypes
```

```
#convert Area from category to numerical
#one hot encoding for three categories
```

```
data = {'Area': ['Urban', 'Suburban', 'Rural', 'Urban', 'Suburban']}
```

```
# Performs one-hot encoding
df_encoded = pd.get_dummies(churn_clean_data, columns=['Area'])
```

```
print(df_encoded)
```

C5.

Please see 'new_cleaned_data_two.csv' file attached to submission

In text citations:

("Completeness", n.d.)

("Importing flat files using pandas", n.d.)

("Is data missing at random?", n.d.)

("Mean, median, & mode imputations", n.d.)

("Imputing using fancyimpute", n.d.)

("Measures of center", n.d.)

("Measures of spread", n.d.)

(Gudikandula, 2018)

Part 4

D1.

Initial linear regression model:

```
from statsmodels.formula.api import ols
```

```
mdl_initial = ols("Tenure ~ Outage_sec_perweek + MonthlyCharge + Area + 0", data =  
churn_clean_data).fit()
```

```
print(mdl_initial.summary())
```

Initial model summary:

```
=====
                        OLS Regression Results
=====
Dep. Variable:          Tenure      R-squared:                0.000
Model:                  OLS        Adj. R-squared:             -0.000
Method:                 Least Squares   F-statistic:              0.5786
Date:                   Tue, 26 Sep 2023   Prob (F-statistic):       0.678
Time:                   14:08:37         Log-Likelihood:          -46732.
No. Observations:      10000           AIC:                    9.347e+04
Df Residuals:          9995            BIC:                    9.351e+04
Df Model:               4
Covariance Type:       nonrobust
=====

```

	coef	std err	t	P> t	[0.025	0.975]
Area[Rural]	35.6743	1.171	30.478	0.000	33.380	37.969
Area[Suburban]	35.0060	1.169	29.948	0.000	32.715	37.297
Area[Urban]	34.9748	1.169	29.914	0.000	32.683	37.267
Outage_sec_perweek	0.0185	0.037	0.498	0.618	-0.054	0.091
MonthlyCharge	-0.0046	0.006	-0.768	0.443	-0.016	0.007

```
=====
Omnibus:                 42080.856   Durbin-Watson:           0.412
Prob(Omnibus):           0.000       Jarque-Bera (JB):        1252.417
Skew:                    0.054       Prob(JB):                1.10e-272
Kurtosis:                1.270       Cond. No.:               1.33e+03
=====
```

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[2] The condition number is large, 1.33e+03. This might indicate that there are strong multicollinearity or other numerical problems.

D2.

To evaluate the performance and accuracy as well as reduce the initial model, I will use the Root Mean Squared Error (RMSE). The RMSE is a robust way to assess the quality of the model's predictions by quantifying the average magnitude of prediction errors. This choice is justified for several reasons:

1. Minimizes prediction errors: To effectively answer the research question, precise and reliable predictions are required. RMSE is designed to measure and minimize prediction errors.
2. Quantifying Model Fit: RMSE quantifies the goodness of fit between the predictive model and the actual data. A lower RMSE signifies that the model's predictions do not vary much from the actual data, ensuring accuracy.
3. Measuring Uncertainty: RMSE measures the spread of residuals and is directly related to the measure of uncertainty. Lower RMSE values indicates a better fitting model, and a higher degree of confidence in the model's predictions. This is essential for providing effective course(s) of action related to the research question.

Utilizing RMSE as the model evaluation metric, will strive to achieve a model that aligns with the research question and demonstrates an accuracy in making predictions. In summary, RMSE will enhance model interpretability, reduce overfitting, and provide precise predictions, which are all necessary pieces of information for addressing the research question effectively.

D3.

Reduced linear regression model:

```
mdl_reduced = ols("Tenure ~ Area_Rural + Area_Suburban + Area_Urban + 0", data =  
feature_churn_data).fit()  
print(mdl_reduced.summary())
```

Reduced model summary:

```
[12]: from statsmodels.formula.api import ols
print(mdl_reduced.summary())
```

```

=====
                        OLS Regression Results
=====
Dep. Variable:          Tenure      R-squared:                0.000
Model:                  OLS        Adj. R-squared:             -0.000
Method:                 Least Squares   F-statistic:              0.7819
Date:                  Fri, 29 Sep 2023   Prob (F-statistic):       0.458
Time:                  20:14:41         Log-Likelihood:          -46732.
No. Observations:      10000          AIC:                     9.347e+04
Df Residuals:          9997          BIC:                     9.349e+04
Df Model:               2
Covariance Type:       nonrobust
=====

```

	coef	std err	t	P> t	[0.025	0.975]
Intercept	25.9684	0.194	133.662	0.000	25.588	26.349
Area_Rural	9.1141	0.372	24.482	0.000	8.384	9.844
Area_Suburban	8.4451	0.372	22.732	0.000	7.717	9.173
Area_Urban	8.4093	0.372	22.589	0.000	7.680	9.139

```

=====
Omnibus:                42066.461   Durbin-Watson:           0.412
Prob(Omnibus):           0.000     Jarque-Bera (JB):        1252.902
Skew:                    0.054     Prob(JB):                 8.63e-273
Kurtosis:                 1.269     Cond. No.                  1.82e+15
=====

```

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[2] The smallest eigenvalue is 4.02e-27. This might indicate that there are strong multicollinearity problems or that the design matrix is singular.

E1.

The data analysis process involved the following:

1. Initial Model Analysis – I initiated the data analysis process by constructing an initial multiple linear regression model. Then, to assess the performance, I computed the RMSE as the model evaluation metric to serve as a base line for comparison.
2. Data Transformation – I recognized that the categorical independent variable, Area, needed to be included into the model, and I employed one-hot encoding. This transformation resulted in the creation of three separate binary variables: Area_Urban, Area_Suburban, and Area_Rural.
3. Feature selection – To refine the model, I choose to conduct recursive feature selection to identify the most influential variables. Through this process, it was determined that the three encoded variables (Area_Urban, Area_Suburban, and Area_Rural) were the most significant contributors to the model's performance.
4. Reduced Model Analysis – Using the selected features, I built a reduced linear regression model. Then, I recalculated the RMSE for the reduced model to gauge its predictive accuracy.
5. Comparison Analysis – Upon comparing the RMSE values of the initial and reduced models, it was determined that both models yielded an RMSE of approximately 25. This outcome indicates that the reduced model, which only includes the most essential features, can achieve a comparable predictive performance with a simpler model.

In text citations:

("More than two explanatory variables", n.d.)

E2.

Output and all calculations of data analysis performed:

1. Initial Model Analysis
 - a. Calculations and output summary provided in D1
 - b. Initial Model RMSE calculation:

```
[10]: residuals = mdl_initial.resid
print(residuals)

0      -27.514666
1      -32.915893
2      -18.672315
3      -17.643798
4      -32.802992
...
9995    33.091200
9996    26.182767
9997   -21.538278
9998    37.066703
9999    29.152857
Length: 10000, dtype: float64

[11]: RSS = np.sum( np.square(residuals) )
print (RSS)

6707904.60854712

[12]: mean_squared_residuals = np.sum( np.square(residuals) ) / len(residuals)
print(mean_squared_residuals)

670.790460854712

[13]: MSE = np.mean( np.square(residuals) )
print(MSE)

670.790460854712

[14]: RMSE = np.sqrt(np.mean( np.square(residuals)))
print(RMSE)

25.899622793676205
```

2. Data Transformation calculations and outputs:

a. One hot encoding:

```
[18]: #Convert Area from category to numerical
#One hot encoding for three categories

data = {'Area': ['Urban', 'Suburban', 'Rural', 'Urban', 'Suburban']}

# Performs one-hot encoding
df_encoded = pd.get_dummies(churn_data, columns=['Area'])
print(df_encoded)
```

```
Unnamed: 0.1 Unnamed: 0.2 CaseOrder Customer_id \
0      0      1      1      1      1      1      1
1      1      2      2      2      2      2      2
2      2      3      3      3      3      3      3
3      3      4      4      4      4      4      4
4      4      5      5      5      5      5      5
...
9995    9995    9995    9995    9995    9995    9995
9996    9996    9997    9997    9997    9997    9997
9997    9997    9998    9998    9998    9998    9998
9998    9998    9999    9999    9999    9999    9999
9999    9999    10000   10000   10000   10000   10000
```

```
Interaction City State \
0  aa9826b0-4161-4a34-ba3e-b8a1c0147776 Point Baker AK
1  f7676459f-c847-4a9d-8a9f-e87d4ac2524 West Branch MI
2  344d214c-3738-4a9d-8a9f-c72c2812d205 Yamhill OR
3  a87a2b4b-3a63-4d96-a13a-89b8c79c111 Del Mar CA
4  68a861f6-8d28-4a51-a587-8a98487e574 Needville TX
...
9995 45d65a2a-aeb4-4518-bf8b-c82d88b8e444 Mount Holly VT
9996 6a985211-4e09-4993-b0da-a1ad411061a Clarksville TN
9997 e838766f-9a81-4fff-3c59-4742e83f624f Hoboken NJ
9998 3775ccfc-8b52-4187-81ae-9057f81ecdf3 Carrollton GA
9999 3775ccfc-8b52-4187-81ae-9057f81ecdf3 Carrollton GA
```

```
County Zip Lat Lng Population \
0 Prince of Wales-Hyder 99927 56.25100 -133.37571 38.0
1 Ogemaw 48661 44.32893 -84.24880 18446.0
2 Yamhill 97148 45.35589 -123.24657 3735.0
3 San Diego 92014 32.96687 -117.24798 13863.0
4 Fort Bend 77461 29.38012 -95.08673 11352.0
...
9995 Rutland 5758 43.43391 -72.78734 648.0
9996 Montgomery 37042 36.56987 -87.41694 77168.0
9997 Wheeler 79061 35.52839 -108.44180 486.0
9998 Carroll 38117 33.58016 -85.13241 35575.0
9999 Habersham 38523 34.70783 -83.53648 12238.0
```

```
Timezone Job Children \
0 America/Sitka Environmental health practitioner 0.005273
1 America/Detroit Programmer, multimedia 1.000000
2 America/Los_Angeles Chief Financial Officer 4.000000
3 America/Los_Angeles Solicitor 1.000000
4 America/Chicago Medical illustrator 0.000000
...
9995 America/New_York Sport and exercise psychologist 3.000000
9996 America/Chicago Consulting civil engineer 4.000000
9997 America/Chicago IT technical support officer 8.873586
9998 America/New_York Water engineer 1.000000
9999 America/New_York Personal assistant 1.000000
```

```
Age Education Employment Income \
0 68.000000 Master's Degree Part Time 28561.990000
1 27.000000 Regular High School Diploma Retired 21784.770000
2 56.000000 Regular High School Diploma Student 21157.834718
3 48.000000 Doctorate Degree Retired 18925.230000
4 83.000000 Master's Degree Student 48074.190000
...
9995 54.872456 Some College, Less than 1 Year Retired 55723.740000
9996 48.000000 Regular High School Diploma Part Time 33267.476603
```

```
DeviceProtection TechSupport StreamingTV StreamingMovies \
0 No No No Yes Yes
1 No No Yes Yes Yes
2 No No No Yes Yes
3 No No Yes No Yes
4 No Yes Yes No No
...
9995 Yes No No No No
9996 Yes No Yes No No
9997 No No No No No
9998 No Yes Yes Yes Yes
9999 Yes No No No No
```

```
PaperlessBilling PaymentMethod Tenure MonthlyCharge \
0 Yes Credit Card (automatic) 6.795513 171.440762
1 Yes Bank Transfer(automatic) 1.156681 242.948015
2 Yes Credit Card (automatic) 15.754144 159.448398
3 Yes Mailed Check 17.487227 128.240493
4 No Mailed Check 1.678972 158.761216
...
9995 No Electronic Check 68.197138 159.828808
9996 No Electronic Check 61.048370 288.856400
9997 Yes Bank Transfer(automatic) 13.446717 168.220908
9998 Yes Credit Card (automatic) 71.095600 252.626600
9999 Yes Electronic Check 63.358860 218.371080
```

```
Bandwidth_GB_Year item1 item2 item3 item4 item5 item6 item7 \
0 904.536110 5 5 5 3 4 4 3
1 800.982766 3 4 3 3 4 3 4
2 2854.780961 4 4 2 4 4 3 3
3 2164.579412 4 4 4 2 5 4 3
4 271.493436 4 4 4 3 4 4 4
...
9995 6511.230808 3 2 3 3 4 3 2
9996 5695.952008 4 5 5 4 4 5 2
```

```
item8 Area_Rural Area_Suburban Area_Urban
0 4 0 0
1 4 0 0
2 3 0 0
3 3 0 1
4 5 0 1
...
9995 3 1 0
9996 5 1 0
9997 5 1 0
9998 4 0 0
9999 1 0 1
```

[10000 rows x 55 columns]

3. Feature selection calculations and outputs:

```
[20]: feature_churn_data = df_encoded[['Tenure', 'Outage_sec_perweek', 'MonthlyCharge',
    'Area_Rural', 'Area_Suburban', 'Area_Urban']]

print(feature_churn_data)
```

	Tenure	Outage_sec_perweek	MonthlyCharge	Area_Rural	Area_Suburban	Area_Urban
0	6.795513	6.972566	171.449762	0	0	1
1	1.156681	12.914541	242.948015	0	0	1
2	15.754144	10.245616	159.448398	0	0	1
3	17.087227	15.286193	120.249493	0	1	0
4	1.678972	8.968316	158.761216	0	0	1
...
9995	68.197130	9.265392	159.828800	1	0	0
9996	61.040378	8.115849	288.856400	1	0	0
9997	13.446717	4.837696	168.220900	1	0	0
9998	71.895600	12.876468	252.628600	0	0	1
9999	63.358860	12.641760	218.371000	0	0	1

```
[10000 rows x 6 columns]
```

```
[21]: #import libraries
from sklearn.feature_selection import RFE
from sklearn.linear_model import LinearRegression

# Extract your target variable (y) and feature matrix (X) from your data
X = feature_churn_data.drop('Tenure', axis=1)
y = feature_churn_data['Tenure']

#estimator (machine learning model) for RFE
estimator = LinearRegression()

# Create the RFE object with the estimator and number of features to select
rfe = RFE(estimator, n_features_to_select=3)
rfe.fit(X, y) # X is the feature matrix, y is the target variable

#fit RFE to the data
rfe.fit(X, y)

#get the selected features and it returns a boolean mask to indicate which features are selected
selected_features = rfe.support_
print("Selected Features:", selected_features)

#rank the selected features and lower rankings mean the features are more important
feature_ranking = rfe.ranking_
print("Feature Ranking:", feature_ranking)

# use the selected features to create a reduced feature matrix
#train your linear regression model with only these features
X_reduced = X.loc[:, selected_features]
print("Reduced Feature Matrix (X_reduced):", X_reduced)

Selected Features: [False False  True  True  True]
Feature Ranking: [2 3 1 1 1]
Reduced Feature Matrix (X_reduced):
```

	Area_Rural	Area_Suburban	Area_Urban
0	0	0	1
1	0	0	1
2	0	0	1
3	0	1	0
4	0	1	0
...
9995	1	0	0
9996	1	0	0
9997	1	0	0
9998	0	0	1
9999	0	0	1

```
[10000 rows x 3 columns]
```

4. Reduced Model Analysis:

a. Calculations and output summary provided in D3

```
[22]: mdl_reduced = ols("Tenure ~ Area_Rural + Area_Suburban + Area_Urban + 0", data = feature_churn_data).fit()
```

```
[26]: from sklearn.model_selection import train_test_split
from sklearn.metrics import mean_squared_error

# Step 3: Split Data into Training and Testing Sets
X = feature_churn_data[['Area_Rural', 'Area_Suburban', 'Area_Urban']]
y = feature_churn_data['Tenure']

# Split data into 80% training and 20% testing
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

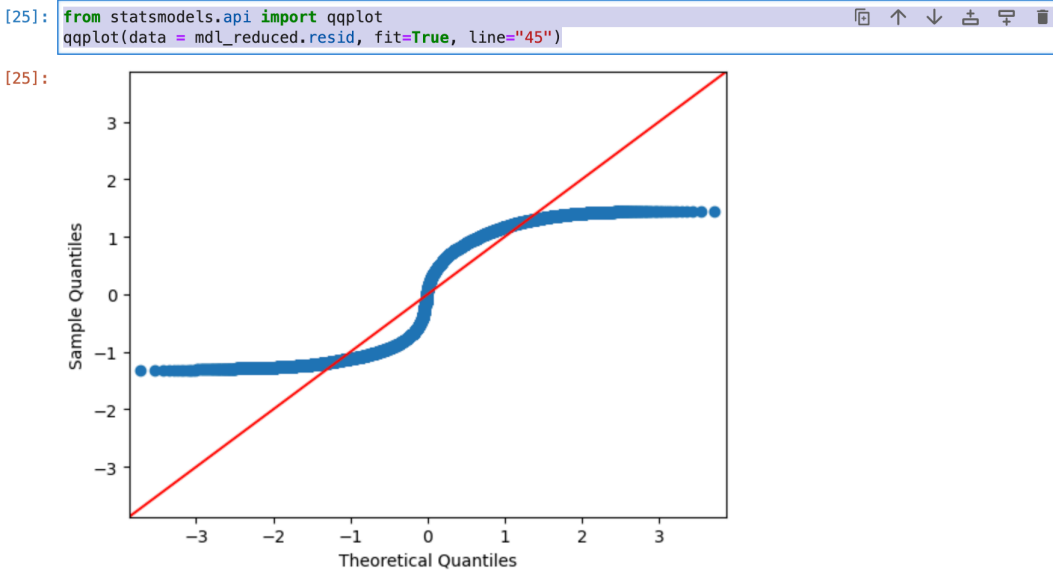
# Step 4: Train the Model
model = LinearRegression()
model.fit(X_train, y_train)

# Step 5: Make Predictions
y_pred = model.predict(X_test)

# Step 6: Evaluate with RMSE
rmse = np.sqrt(mean_squared_error(y_test, y_pred))
print("Root Mean Squared Error (RMSE):", rmse)
```

Root Mean Squared Error (RMSE): 25.923698189470592

5. Residual plot for the reduced model (qqplot):



6. Model's residual standard error for the reduced model:

```
[30]: #residual standard error of reduced model
      mse = mdl_reduced.mse_resid
      rse = np.sqrt(mse)
      print('rse:', rse)

rse: 25.904481450646614
```

E3. Please see Churn Clean Data.ipynb for code

In text citations:

("Goodness-of-Fit", n.d.)

("The curse of dimensionality", n.d.)

("Selecting features for model performance", n.d.)

(Lekhana_Ganji, 2023)

("Working with model objects", n.d.)

Part 5

F1.

Data Summary:

Regression equation for the reduced model:

$$Y = 9.1141 \cdot \text{Area_Rural} + 8.4451 \cdot \text{Area_Suburban} + 8.4093 \cdot \text{Area_Urban} + 25.9684$$

Interpretation of the coefficients of the reduced model:

Y = the dependent variable the model is predicting

β_0 = The constant intercept term

β_1 = the coefficient for the Area Rural independent variable

β_2 = the coefficient for the Area Suburban independent variable

β_3 = the coefficient for the Area Urban independent variable

The independent variable coefficients represent the change in the dependent variable for one unit change in the corresponding independent variable, while holding all other variables constant.

Statistical and practical significance of the reduced model:

The p-values for the independent variables were below a well-known significance level of $\alpha = 0.05$, and therefore, were found to be statistically significant. This indicates that these independent variables have a statistically significant impact on the dependent variable.

While the coefficients of the independent variables are statistically significant, it is also important to consider their practical significance. The reduced model could be meaningful in the real world if telecommunication companies want to assess where the highest customer churn occurs based on geographical location.

Data Analysis limitations:

While the analysis has produced meaningful insights, the analysis process limits the telecommunication companies from determining other predictor factors for customer churn. The initial model, which serves as the foundation for feature selection and reduction was constructed with a limited set of independent variables to reduce overall model complexity. This constraint limited the range of independent variables available for consideration during feature selection.

After the recursive feature selection process, the selected features were focused on the customer's geographic location as a significant predictor of customer churn can be a limitation. Although meaningful, it limits the companies' ability to explore and identify other potential predictor factors that might influence churn behavior. In conclusion, customer churn can be a complex decision that can be influenced by a plethora of factors, and more independent variables could have helped to relieve some limitations.

F2.

Based on my results, the recommended course of action would be to investigate how customer's geographic location affects customer's churn. Using customer's geographical location, the telecommunication companies should zone in on one geographical location, and then rerun the initial model with additional independent variables, and the predicted dependent variable being the geographical location where customer churn is highly likely. By using this course of action, the telecommunications companies can develop detailed insights to proactively prevent customer churn.

In text citations:

(LaMorte, 2016)

("Inference on transformed variables", n.d.)

Part 6

H.

Citations for code:

DataCamp. (n.d.). More than two explanatory variables [Video file]. Retrieved from <https://campus.datacamp.com/courses/intermediate-regression-with-statsmodels-in-python/multiple-linear-regression-3?ex=7>

DataCamp. (n.d.). Completeness [Video file]. Retrieved from <https://campus.datacamp.com/courses/cleaning-data-in-python/advanced-data-problems-3?ex=8>

DataCamp. (n.d.). Importing flat files using pandas [Video file]. Retrieved from <https://campus.datacamp.com/courses/introduction-to-importing-data-in-python/introduction-and-flat-files-1?ex=15>

DataCamp. (n.d.). Mean, median, & mode imputations [Video file]. Retrieved from <https://campus.datacamp.com/courses/dealing-with-missing-data-in-python/imputation-techniques?ex=1>

DataCamp. (n.d.). Imputing using fancyimpute [Video file]. Retrieved from <https://campus.datacamp.com/courses/dealing-with-missing-data-in-python/advanced-imputation-techniques?ex=1>

DataCamp. (n.d.). The curse of dimensionality [Video file]. Retrieved from <https://campus.datacamp.com/courses/dimensionality-reduction-in-python/feature-selection-i-selecting-for-feature-information?ex=1>

DataCamp. (n.d.). Selecting features for model performance [Video file]. Retrieved from <https://campus.datacamp.com/courses/dimensionality-reduction-in-python/feature-selection-ii-selecting-for-model-accuracy?ex=1>

Lekhana_Ganji. (2023, April 18). Machine Learning - One Hot Encoding of Datasets in Python. GeeksforGeeks. <https://www.geeksforgeeks.org/ml-one-hot-encoding-of-datasets-in-python/>

DataCamp. (n.d.). Working with model objects [Video file]. Retrieved from <https://campus.datacamp.com/courses/introduction-to-regression-with-statsmodels-in-python/predictions-and-model-objects-2?ex=5>

DataCamp. (n.d.). Visualizing model fit [Video file]. Retrieved from <https://campus.datacamp.com/courses/introduction-to-regression-with-statsmodels-in-python/assessing-model-fit-e78fd9fe-6303-4048-8748-33b19c4222fe?ex=4>

DataCamp. (n.d.). Quantifying model fit [Video file]. Retrieved from <https://campus.datacamp.com/courses/introduction-to-regression-with-statsmodels-in-python/assessing-model-fit-e78fd9fe-6303-4048-8748-33b19c4222fe?ex=1>

I.

Citations for content:

DataCamp. (n.d.). Technical conditions for linear regression [Video file]. Retrieved from <https://campus.datacamp.com/courses/inference-for-linear-regression-in-r/technical-conditions-in-linear-regression?ex=1>

DataCamp. (n.d.). Goodness-of-fit [Video file]. Retrieved from <https://campus.datacamp.com/courses/introduction-to-linear-modeling-in-python/making-model-predictions?ex=8>

DataCamp. (n.d.). Is the data missing at random? [Video file]. Retrieved from <https://campus.datacamp.com/courses/dealing-with-missing-data-in-python/does-missingness-have-a-pattern?ex=1>

DataCamp. (n.d.). Measures of center [Video file]. Retrieved from <https://campus.datacamp.com/courses/introduction-to-statistics-in-python/summary-statistics-1?ex=4>

DataCamp. (n.d.). Measures of spread [Video file]. Retrieved from <https://campus.datacamp.com/courses/introduction-to-statistics-in-python/summary-statistics-1?ex=7>

DataCamp. (n.d.). Inference on transformed variables [Video file]. Retrieved from <https://campus.datacamp.com/courses/inference-for-linear-regression-in-r/building-on-inference-in-simple-linear-regression?ex=1>

LaMorte, W. (2016, May 31). The Multiple Linear Regression Equation. Multivariate Methods. https://sphweb.bumc.bu.edu/otlt/mph-modules/bs/bs704-ep713_multivariablemethods/bs704-ep713_multivariablemethods2.html

Gudikandula, P. (2018, November 9). Exploratory Data Analysis (beginner), Univariate, Bivariate and Multivariate – Habberman dataset. [purnasaigudikandula.medium.com](https://purnasaigudikandula.medium.com/exploratory-data-analysis-beginner-univariate-bivariate-and-multivariate-habberman-dataset-2365264b751). Retrieved from <https://purnasaigudikandula.medium.com/exploratory-data-analysis-beginner-univariate-bivariate-and-multivariate-habberman-dataset-2365264b751>